

1. Apply data cleaning techniques on any dataset (e.g., Paper Reviews dataset in UCI repository). Techniques may include handling missing values, outliers and inconsistent values. A set of validation rules can be prepared based on the dataset and validations can be performed.

```
In [34]: import pandas as pd
import json

with open('reviews.json', 'r', encoding='utf-8') as f:
    data = json.load(f)

df = pd.json_normalize(
    data['paper'],
    record_path=['review'],
    meta=['id', 'preliminary_decision'],
    record_prefix='review_',
    meta_prefix='paper_'
)

print(df.head())
```

|   | review_confidence | review_evaluation | review_id | review_lang |  |
|---|-------------------|-------------------|-----------|-------------|--|
| 0 | 4                 | 1                 | 1         | es          |  |
| 1 | 4                 | 1                 | 2         | es          |  |
| 2 | 5                 | 1                 | 3         | es          |  |
| 3 | 4                 | 2                 | 1         | es          |  |
| 4 | 4                 | 2                 | 2         | es          |  |

|   | review_orientation | review_remarks |  |
|---|--------------------|----------------|--|
| 0 | 0                  |                |  |
| 1 | 1                  |                |  |
| 2 | 1                  |                |  |
| 3 | 1                  |                |  |
| 4 | 0                  |                |  |

|   | review_text                                       | review_timespan | paper_id |  |
|---|---------------------------------------------------|-----------------|----------|--|
| 0 | - El artículo aborda un problema contingente y... | 2010-07-05      | 1        |  |
| 1 | El artículo presenta recomendaciones prácticas... | 2010-07-05      | 1        |  |
| 2 | - El tema es muy interesante y puede ser de mu... | 2010-07-05      | 1        |  |
| 3 | Se explica en forma ordenada y didáctica una e... | 2010-07-05      | 2        |  |
| 4 |                                                   | 2010-07-05      | 2        |  |

|   | paper_preliminary_decision |
|---|----------------------------|
| 0 | accept                     |
| 1 | accept                     |
| 2 | accept                     |
| 3 | accept                     |
| 4 | accept                     |

```
In [35]: # ===== handling missing values =====

print(df.isnull().sum())
```

```

df['review_confidence'] = pd.to_numeric(df['review_confidence'], errors='coerce')
df['review_evaluation'] = pd.to_numeric(df['review_evaluation'], errors='coerce')

df.fillna({'review_confidence': df['review_confidence'].mean(), 'review_evaluation':

# df['review_remarks'].fillna('No remarks', inplace=True)
# df['review_text'].fillna('No text provided', inplace=True)

df.fillna({'review_remarks': '', 'review_text': ''}, inplace=True)

review_confidence      2
review_evaluation      0
review_id              0
review_lan              0
review_orientation      0
review_remarks          0
review_text            0
review_timespan         0
paper_id               0
paper_preliminary_decision  0
dtype: int64

```

```

In [36]: # ===== handling inconsistent data =====

df['review_lan'] = df['review_lan'].str.strip().str.lower()
# df['review_lan'].replace({'spa': 'es', 'en-us': 'en'}, inplace=True)
df.replace({'review_lan': {'spa': 'es', 'en-us': 'en'}}, inplace=True)

df['paper_preliminary_decision'] = df['paper_preliminary_decision'].str.strip().str
df.loc[~df['paper_preliminary_decision'].isin(['accept', 'reject']), 'paper_prelimi

```

```

In [37]: # Visualize outliers before handling
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
df.boxplot(column=['review_confidence'])
plt.title('Before: Review Confidence')

plt.subplot(1, 2, 2)
df.boxplot(column=['review_evaluation'])
plt.title('Before: Review Evaluation')
plt.tight_layout()
plt.show()

```



```
In [38]: # ===== handling outliers using IQR method =====

for col in ['review_confidence', 'review_evaluation']:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    outliers = df[(df[col] < lower_bound) | (df[col] > upper_bound)]
    print(f"Outliers found in {col}: {len(outliers)}")

    # Remove outliers
    df = df[(df[col] >= lower_bound) & (df[col] <= upper_bound)]

print(f"\nShape of DataFrame after removing outliers: {df.shape}")
```


Outliers found in review\_confidence: 7

Outliers found in review\_evaluation: 0

Shape of DataFrame after removing outliers: (398, 10)

```
In [39]: # Visualize data after handling outliers
plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
df.boxplot(column=['review_confidence'])
plt.title('After: Review Confidence')

plt.subplot(1, 2, 2)
df.boxplot(column=['review_evaluation'])
plt.title('After: Review Evaluation')
plt.tight_layout()
plt.show()
```



Validation Rules:

| Rule | Field                      | Expected Range / Type     |
|------|----------------------------|---------------------------|
| R1   | review_confidence          | 1–5                       |
| R2   | review_evaluation          | -2–2                      |
| R3   | review_lan                 | one of ['en', 'es', 'fr'] |
| R4   | paper_preliminary_decision | ['accept', 'reject']      |
| R5   | review_text                | not empty                 |

```
In [42]: # ===== applying validation rules =====

rule1 = df["review_confidence"].between(1, 5)
rule2 = df["review_evaluation"].between(-2, 2)
rule3 = df["review_lan"].isin(["en", "es", "fr"])
rule4 = df["paper_preliminary_decision"].isin(["accept", "reject"])
rule5 = df["review_text"].str.strip().ne("")

Rules = pd.DataFrame({
    "Rule 1 (Confidence 1-5)": rule1,
    "Rule 2 (Evaluation -2-2)": rule2,
    "Rule 3 (Language Valid)": rule3,
    "Rule 4 (Decision Valid)": rule4,
    "Rule 5 (Text Non-empty)": rule5
})

print("Rule Violation Summary:")
for col in Rules.columns:
    total = len(Rules)
    passed = Rules[col].sum()
    failed = total - passed
    print(f"{col}: {failed} violations out of {total}")
```

Rule Violation Summary:

Rule 1 (Confidence 1-5): 0 violations out of 398

Rule 2 (Evaluation -2-2): 0 violations out of 398

Rule 3 (Language Valid): 0 violations out of 398

Rule 4 (Decision Valid): 21 violations out of 398

Rule 5 (Text Non-empty): 6 violations out of 398

```
In [41]: # ===== final cleaned data =====

print(df.describe(include='all'))

# Visualize numeric distributions
import matplotlib.pyplot as plt
df['review_confidence'].hist()
plt.title('Distribution of Review Confidence')
plt.show()

df['review_evaluation'].hist()
plt.title('Distribution of Review Evaluation')
plt.show()


# save cleaned data
df.to_csv("cleaned_paper_reviews.csv", index=False)
```

|        | review_confidence | review_evaluation | review_id  | review_lang | \ |
|--------|-------------------|-------------------|------------|-------------|---|
| count  | 398.000000        | 398.000000        | 398.000000 | 398         |   |
| unique | NaN               | NaN               | NaN        | 2           |   |
| top    | NaN               | NaN               | NaN        | es          |   |
| freq   | NaN               | NaN               | NaN        | 381         |   |
| mean   | 3.618458          | 0.188442          | 1.819095   | NaN         |   |
| std    | 0.776587          | 1.504709          | 0.816972   | NaN         |   |
| min    | 2.000000          | -2.000000         | 1.000000   | NaN         |   |
| 25%    | 3.000000          | -1.000000         | 1.000000   | NaN         |   |
| 50%    | 4.000000          | 0.000000          | 2.000000   | NaN         |   |
| 75%    | 4.000000          | 2.000000          | 2.000000   | NaN         |   |
| max    | 5.000000          | 2.000000          | 4.000000   | NaN         |   |


|        | review_orientation | review_remarks | review_text | review_timespan | \ |
|--------|--------------------|----------------|-------------|-----------------|---|
| count  | 398                | 398            | 398         | 398             |   |
| unique | 5                  | 108            | 393         | 4               |   |
| top    | -1                 |                |             | 2014-07-05      |   |
| freq   | 139                | 288            | 6           | 124             |   |
| mean   | NaN                | NaN            | NaN         | NaN             |   |
| std    | NaN                | NaN            | NaN         | NaN             |   |
| min    | NaN                | NaN            | NaN         | NaN             |   |
| 25%    | NaN                | NaN            | NaN         | NaN             |   |
| 50%    | NaN                | NaN            | NaN         | NaN             |   |
| 75%    | NaN                | NaN            | NaN         | NaN             |   |
| max    | NaN                | NaN            | NaN         | NaN             |   |

|        | paper_id | paper_preliminary_decision |
|--------|----------|----------------------------|
| count  | 398.0    | 377                        |
| unique | 168.0    | 2                          |
| top    | 112.0    | accept                     |
| freq   | 4.0      | 257                        |
| mean   | NaN      | NaN                        |
| std    | NaN      | NaN                        |
| min    | NaN      | NaN                        |
| 25%    | NaN      | NaN                        |
| 50%    | NaN      | NaN                        |
| 75%    | NaN      | NaN                        |
| max    | NaN      | NaN                        |

Distribution of Review Confidence



Distribution of Review Evaluation



## 2. Apply data pre-processing techniques such as standardization/normalization, transformation, aggregation, discretization/binarization, sampling etc. on any dataset

```
In [1]: import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler, MinMaxScaler, Binarizer

iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df['species'] = pd.Categorical.from_codes(iris.target, iris.target_names)

print(df.head())
print("\nDataset shape:", df.shape)
print("\nSummary statistics:\n", df.describe())

X = df.drop(columns=['species'])
```



|   | sepal length (cm) | sepal width (cm) | petal length (cm) | petal width (cm) | \ |
|---|-------------------|------------------|-------------------|------------------|---|
| 0 | 5.1               | 3.5              | 1.4               | 0.2              |   |
| 1 | 4.9               | 3.0              | 1.4               | 0.2              |   |
| 2 | 4.7               | 3.2              | 1.3               | 0.2              |   |
| 3 | 4.6               | 3.1              | 1.5               | 0.2              |   |
| 4 | 5.0               | 3.6              | 1.4               | 0.2              |   |

|   | species |
|---|---------|
| 0 | setosa  |
| 1 | setosa  |
| 2 | setosa  |
| 3 | setosa  |
| 4 | setosa  |

Dataset shape: (150, 5)

Summary statistics:

|       | sepal length (cm) | sepal width (cm) | petal length (cm) | \ |
|-------|-------------------|------------------|-------------------|---|
| count | 150.000000        | 150.000000       | 150.000000        |   |
| mean  | 5.843333          | 3.057333         | 3.758000          |   |
| std   | 0.828066          | 0.435866         | 1.765298          |   |
| min   | 4.300000          | 2.000000         | 1.000000          |   |
| 25%   | 5.100000          | 2.800000         | 1.600000          |   |
| 50%   | 5.800000          | 3.000000         | 4.350000          |   |
| 75%   | 6.400000          | 3.300000         | 5.100000          |   |
| max   | 7.900000          | 4.400000         | 6.900000          |   |

|       | petal width (cm) |
|-------|------------------|
| count | 150.000000       |
| mean  | 1.199333         |
| std   | 0.762238         |
| min   | 0.100000         |
| 25%   | 0.300000         |
| 50%   | 1.300000         |
| 75%   | 1.800000         |
| max   | 2.500000         |

```
In [2]: scaler = StandardScaler()
df_standardized = scaler.fit_transform(X)
df_standardized = pd.DataFrame(df_standardized, columns=X.columns)

print("\nStandardized Data (first 5 rows):\n", df_standardized.head())
```

Standardized Data (first 5 rows):

|   | sepal length (cm) | sepal width (cm) | petal length (cm) | petal width (cm) |
|---|-------------------|------------------|-------------------|------------------|
| 0 | -0.900681         | 1.019004         | -1.340227         | -1.315444        |
| 1 | -1.143017         | -0.131979        | -1.340227         | -1.315444        |
| 2 | -1.385353         | 0.328414         | -1.397064         | -1.315444        |
| 3 | -1.506521         | 0.098217         | -1.283389         | -1.315444        |
| 4 | -1.021849         | 1.249201         | -1.340227         | -1.315444        |

```
In [3]: normalizer = MinMaxScaler()
df_normalized = normalizer.fit_transform(X)
df_normalized = pd.DataFrame(df_normalized, columns=X.columns)

print("\nNormalized Data (first 5 rows):\n", df_normalized.head())
```

Normalized Data (first 5 rows):

|   | sepal length (cm) | sepal width (cm) | petal length (cm) | petal width (cm) |
|---|-------------------|------------------|-------------------|------------------|
| 0 | 0.222222          | 0.625000         | 0.067797          | 0.041667         |
| 1 | 0.166667          | 0.416667         | 0.067797          | 0.041667         |
| 2 | 0.111111          | 0.500000         | 0.050847          | 0.041667         |
| 3 | 0.083333          | 0.458333         | 0.084746          | 0.041667         |
| 4 | 0.194444          | 0.666667         | 0.067797          | 0.041667         |

```
In [4]: df_transformed = df.copy()
df_transformed["sepal width (cm)"] = np.log1p(df_transformed["sepal width (cm)"])

print("\nTransformed (log) Data (first 5 rows):\n", df_transformed.head())
```

Transformed (log) Data (first 5 rows):

|   | sepal length (cm) | sepal width (cm) | petal length (cm) | petal width (cm) | \ |
|---|-------------------|------------------|-------------------|------------------|---|
| 0 | 5.1               | 1.504077         | 1.4               | 0.2              |   |
| 1 | 4.9               | 1.386294         | 1.4               | 0.2              |   |
| 2 | 4.7               | 1.435085         | 1.3               | 0.2              |   |
| 3 | 4.6               | 1.410987         | 1.5               | 0.2              |   |
| 4 | 5.0               | 1.526056         | 1.4               | 0.2              |   |

|   | species |
|---|---------|
| 0 | setosa  |
| 1 | setosa  |
| 2 | setosa  |
| 3 | setosa  |
| 4 | setosa  |

```
In [5]: agg_df = df.groupby("species", observed=False).mean()
print("\nAggregated Mean by Species:\n", agg_df)
```

Aggregated Mean by Species:

|            | sepal length (cm) | sepal width (cm) | petal length (cm) | \ |
|------------|-------------------|------------------|-------------------|---|
| species    |                   |                  |                   |   |
| setosa     | 5.006             | 3.428            | 1.462             |   |
| versicolor | 5.936             | 2.770            | 4.260             |   |
| virginica  | 6.588             | 2.974            | 5.552             |   |

|            | petal width (cm) |
|------------|------------------|
| species    |                  |
| setosa     | 0.246            |
| versicolor | 1.326            |
| virginica  | 2.026            |

```
In [6]: bins = [0, 2.5, 5, 7.5]
labels = ['Small', 'Medium', 'Large']

df_discrete = df.copy()
df_discrete['custom_petal_length'] = pd.cut(df['petal length (cm)'], bins=bins, labels=labels)

print("\nDiscretized Data (Petal Length into 3 bins):\n", df_discrete.head())
```

Discretized Data (Petal Length into 3 bins):

|   | sepal length (cm) | sepal width (cm) | petal length (cm) | petal width (cm) | \ |
|---|-------------------|------------------|-------------------|------------------|---|
| 0 | 5.1               | 3.5              | 1.4               | 0.2              |   |
| 1 | 4.9               | 3.0              | 1.4               | 0.2              |   |
| 2 | 4.7               | 3.2              | 1.3               | 0.2              |   |
| 3 | 4.6               | 3.1              | 1.5               | 0.2              |   |
| 4 | 5.0               | 3.6              | 1.4               | 0.2              |   |

|   | species | custom_petal_length |
|---|---------|---------------------|
| 0 | setosa  | Small               |
| 1 | setosa  | Small               |
| 2 | setosa  | Small               |
| 3 | setosa  | Small               |
| 4 | setosa  | Small               |

```
In [7]: binarizer = Binarizer(threshold=1.0)
df_binarized = df.copy()
df_binarized["wide_petal"] = binarizer.fit_transform(df[["petal width (cm)"]])

print("\nBinarized Data (Petal width >= 1.0 -> 1):\n", df_binarized.head())
```

Binarized Data (Petal width >= 1.0 -> 1):

|   | sepal length (cm) | sepal width (cm) | petal length (cm) | petal width (cm) | \ |
|---|-------------------|------------------|-------------------|------------------|---|
| 0 | 5.1               | 3.5              | 1.4               | 0.2              |   |
| 1 | 4.9               | 3.0              | 1.4               | 0.2              |   |
| 2 | 4.7               | 3.2              | 1.3               | 0.2              |   |
| 3 | 4.6               | 3.1              | 1.5               | 0.2              |   |
| 4 | 5.0               | 3.6              | 1.4               | 0.2              |   |

|   | species | wide_petal |
|---|---------|------------|
| 0 | setosa  | 0.0        |
| 1 | setosa  | 0.0        |
| 2 | setosa  | 0.0        |
| 3 | setosa  | 0.0        |
| 4 | setosa  | 0.0        |

```
In [8]: sample_df = df.sample(frac=0.2, random_state=42)
print("\nSampled 20% of the Data:\n", sample_df)
```

Sampled 20% of the Data:

|     | sepal length (cm) | sepal width (cm) | petal length (cm) | petal width (cm) | \ |
|-----|-------------------|------------------|-------------------|------------------|---|
| 73  | 6.1               | 2.8              | 4.7               | 1.2              |   |
| 18  | 5.7               | 3.8              | 1.7               | 0.3              |   |
| 118 | 7.7               | 2.6              | 6.9               | 2.3              |   |
| 78  | 6.0               | 2.9              | 4.5               | 1.5              |   |
| 76  | 6.8               | 2.8              | 4.8               | 1.4              |   |
| 31  | 5.4               | 3.4              | 1.5               | 0.4              |   |
| 64  | 5.6               | 2.9              | 3.6               | 1.3              |   |
| 141 | 6.9               | 3.1              | 5.1               | 2.3              |   |
| 68  | 6.2               | 2.2              | 4.5               | 1.5              |   |
| 82  | 5.8               | 2.7              | 3.9               | 1.2              |   |
| 110 | 6.5               | 3.2              | 5.1               | 2.0              |   |
| 12  | 4.8               | 3.0              | 1.4               | 0.1              |   |
| 36  | 5.5               | 3.5              | 1.3               | 0.2              |   |
| 9   | 4.9               | 3.1              | 1.5               | 0.1              |   |
| 19  | 5.1               | 3.8              | 1.5               | 0.3              |   |
| 56  | 6.3               | 3.3              | 4.7               | 1.6              |   |
| 104 | 6.5               | 3.0              | 5.8               | 2.2              |   |
| 69  | 5.6               | 2.5              | 3.9               | 1.1              |   |
| 55  | 5.7               | 2.8              | 4.5               | 1.3              |   |
| 132 | 6.4               | 2.8              | 5.6               | 2.2              |   |
| 29  | 4.7               | 3.2              | 1.6               | 0.2              |   |
| 127 | 6.1               | 3.0              | 4.9               | 1.8              |   |
| 26  | 5.0               | 3.4              | 1.6               | 0.4              |   |
| 128 | 6.4               | 2.8              | 5.6               | 2.1              |   |
| 131 | 7.9               | 3.8              | 6.4               | 2.0              |   |
| 145 | 6.7               | 3.0              | 5.2               | 2.3              |   |
| 108 | 6.7               | 2.5              | 5.8               | 1.8              |   |
| 143 | 6.8               | 3.2              | 5.9               | 2.3              |   |
| 45  | 4.8               | 3.0              | 1.4               | 0.3              |   |
| 30  | 4.8               | 3.1              | 1.6               | 0.2              |   |

|     | species    |
|-----|------------|
| 73  | versicolor |
| 18  | setosa     |
| 118 | virginica  |
| 78  | versicolor |
| 76  | versicolor |
| 31  | setosa     |
| 64  | versicolor |
| 141 | virginica  |
| 68  | versicolor |
| 82  | versicolor |
| 110 | virginica  |
| 12  | setosa     |
| 36  | setosa     |
| 9   | setosa     |
| 19  | setosa     |
| 56  | versicolor |
| 104 | virginica  |
| 69  | versicolor |
| 55  | versicolor |
| 132 | virginica  |
| 29  | setosa     |
| 127 | virginica  |

|     |           |
|-----|-----------|
| 26  | setosa    |
| 128 | virginica |
| 131 | virginica |
| 145 | virginica |
| 108 | virginica |
| 143 | virginica |
| 45  | setosa    |
| 30  | setosa    |

### 3. Run Apriori algorithm to find frequent item sets and association rules on 2 real datasets and use appropriate evaluation measures to compute correctness of obtained patterns

- a) Use minimum support as 50% and minimum confidence as 75%
- b) Use minimum support as 60% and minimum confidence as 60 %

```
In [1]: from efficient_apriori import apriori
import pandas as pd
import numpy as np

groceries = pd.read_csv("groceries.csv", header=0)
print(groceries.head())
print(groceries.shape)

bakery = pd.read_csv("bakery.csv")
bakery.columns = bakery.columns.str.strip()
print(bakery.head())
print(bakery.shape)
```

|   | Item(s) | Item 1           | Item 2              | Item 3 \       |
|---|---------|------------------|---------------------|----------------|
| 0 | 4       | citrus fruit     | semi-finished bread | margarine      |
| 1 | 3       | tropical fruit   | yogurt              | coffee         |
| 2 | 1       | whole milk       | NaN                 | NaN            |
| 3 | 4       | pip fruit        | yogurt              | cream cheese   |
| 4 | 4       | other vegetables | whole milk          | condensed milk |

|   |           | Item 4         | Item 5 | Item 6 | Item 7 | Item 8 | Item 9 | ... Item 23 \ |
|---|-----------|----------------|--------|--------|--------|--------|--------|---------------|
| 0 |           | ready soups    | NaN    | NaN    | NaN    | NaN    | NaN    | NaN           |
| 1 |           | NaN            | NaN    | NaN    | NaN    | NaN    | NaN    | NaN           |
| 2 |           | NaN            | NaN    | NaN    | NaN    | NaN    | NaN    | NaN           |
| 3 |           | meat spreads   | NaN    | NaN    | NaN    | NaN    | NaN    | NaN           |
| 4 | long life | bakery product | NaN    | NaN    | NaN    | NaN    | NaN    | NaN           |

|   | Item 24 | Item 25 | Item 26 | Item 27 | Item 28 | Item 29 | Item 30 | Item 31 | Item 32 |
|---|---------|---------|---------|---------|---------|---------|---------|---------|---------|
| 0 | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     |
| 1 | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     |
| 2 | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     |
| 3 | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     |
| 4 | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     | NaN     |

[5 rows x 33 columns]  
(9835, 33)

|   | ID | Chocolate Cake | Lemon Cake | Casino Cake | Opera Cake | Strawberry Cake \ |
|---|----|----------------|------------|-------------|------------|-------------------|
| 0 | 1  | 0              | 0          | 0           | 0          | 1                 |
| 1 | 2  | 0              | 1          | 0           | 0          | 0                 |
| 2 | 3  | 0              | 0          | 0           | 0          | 0                 |
| 3 | 4  | 0              | 0          | 0           | 0          | 0                 |
| 4 | 5  | 0              | 0          | 0           | 0          | 0                 |

|   | Truffle Cake | Chocolate Eclair | Coffee Eclair | Vanilla Eclair | ... \ |
|---|--------------|------------------|---------------|----------------|-------|
| 0 | 1            | 1                | 0             | 0              | ...   |
| 1 | 0            | 0                | 0             | 0              | ...   |
| 2 | 0            | 0                | 0             | 0              | ...   |
| 3 | 0            | 0                | 0             | 0              | ...   |
| 4 | 0            | 0                | 0             | 0              | ...   |

|   | Lemon Lemonade | Raspberry Lemonade | Orange Juice | Green Tea | Bottled Water \ |
|---|----------------|--------------------|--------------|-----------|-----------------|
| 0 | 0              | 0                  | 0            | 0         | 0               |
| 1 | 0              | 0                  | 0            | 0         | 0               |
| 2 | 0              | 0                  | 0            | 0         | 0               |
| 3 | 0              | 0                  | 0            | 0         | 0               |
| 4 | 0              | 0                  | 0            | 0         | 0               |

|   | Hot Coffee | Chocolate Coffee | Vanilla Frappuccino | Cherry Soda \ |
|---|------------|------------------|---------------------|---------------|
| 0 | 0          | 0                | 0                   | 0             |
| 1 | 0          | 1                | 1                   | 0             |
| 2 | 0          | 0                | 0                   | 0             |
| 3 | 0          | 0                | 0                   | 0             |
| 4 | 0          | 0                | 0                   | 0             |

|   | Single Espresso |
|---|-----------------|
| 0 | 0               |
| 1 | 0               |
| 2 | 0               |
| 3 | 0               |

[5 rows x 51 columns]  
(5000, 51)

```
In [2]: transactions = []
        for i in range(len(groceries)):
            row = groceries.iloc[i, 1:].dropna().astype(str).tolist()
            transactions.append(row)

        print(len(transactions), transactions[:5])

        transactions_b = []
        item_columns = bakery.columns[1:]
        for i in range(len(bakery)):
            items = item_columns[bakery.iloc[i, 1:] == 1].tolist()
            if items:
                transactions_b.append(items)

        print(len(transactions_b), transactions_b[:5])
```

9835 [['citrus fruit', 'semi-finished bread', 'margarine', 'ready soups'], ['tropical fruit', 'yogurt', 'coffee'], ['whole milk'], ['pip fruit', 'yogurt', 'cream cheese', 'meat spreads'], ['other vegetables', 'whole milk', 'condensed milk', 'long life bakery product']]

5000 [['Strawberry Cake', 'Truffle Cake', 'Chocolate Eclair', 'Almond Tart'], ['Lemon Cake', 'Apple Tart', 'Lemon Tart', 'Chocolate Coffee', 'Vanilla Frappuccino'], ['Blueberry Tart', 'Apricot Croissant'], ['Cherry Tart', 'Lemon Cookie', 'Apricot Danish'], ['Apple Tart', 'Gongolais Cookie', 'Marzipan Cookie', 'Almond Croissant']]

```
In [3]: itemsets, rules = apriori(
        transactions,
        min_support=0.05,
        min_confidence=0.25
    )

    print("Frequent Itemsets (Groceries):")
    print(itemsets)
    print("\nRules (Groceries):")
    print(rules)

    b_item, b_rules = apriori(
        transactions_b,
        min_support=0.05,
        min_confidence=0.25
    )

    print("\nFrequent Itemsets (Bakery):")
    print(b_item)
    print("\nRules (Bakery):")
    print(b_rules)
```



#### Frequent Itemsets (Groceries):

{1: {'citrus fruit',): 814, ('margarine',): 576, ('tropical fruit',): 1032, ('yogurt',): 1372, ('coffee',): 571, ('whole milk',): 2513, ('pip fruit',): 744, ('other vegetables',): 1903, ('butter',): 545, ('rolls/buns',): 1809, ('bottled beer',): 792, ('bottled water',): 1087, ('curd',): 524, ('beef',): 516, ('frankfurter',): 580, ('soda',): 1715, ('fruit/vegetable juice',): 711, ('newspapers',): 785, ('pastry',): 875, ('root vegetables',): 1072, ('canned beer',): 764, ('sausage',): 924, ('brown bread',): 638, ('shopping bags',): 969, ('napkins',): 515, ('pork',): 567, ('whipped/sour cream',): 705, ('domestic eggs',): 624}, 2: {'other vegetables', 'whole milk': 736, ('rolls/buns', 'whole milk': 557, ('whole milk', 'yogurt'): 551}}

#### Rules (Groceries):

[{whole milk} -> {other vegetables}, {other vegetables} -> {whole milk}, {rolls/buns} -> {whole milk}, {yogurt} -> {whole milk}]

#### Frequent Itemsets (Bakery):

{1: {'Strawberry Cake',): 480, ('Truffle Cake',): 438, ('Lemon Cake',): 425, ('Apple Tart',): 367, ('Lemon Tart',): 354, ('Chocolate Coffee',): 404, ('Vanilla Frappuccino',): 367, ('Blueberry Tart',): 426, ('Apricot Croissant',): 411, ('Cherry Tart',): 460, ('Lemon Cookie',): 321, ('Apricot Danish',): 446, ('Gongolais Cookie',): 477, ('Marzipan Cookie',): 418, ('Walnut Cookie',): 353, ('Chocolate Cake',): 399, ('Casino Cake',): 373, ('Berry Tart',): 391, ('Lemon Lemonade',): 324, ('Bottled Water',): 386, ('Opera Cake',): 418, ('Tuile Cookie',): 499, ('Apple Croissant',): 371, ('Apple Danish',): 391, ('Cherry Soda',): 335, ('Chocolate Tart',): 381, ('Coffee Eclair',): 554, ('Apple Pie',): 391, ('Almond Twist',): 407, ('Orange Juice',): 461, ('Raspberry Lemonade',): 339, ('Napolean Cake',): 409, ('Hot Coffee',): 513, ('Raspberry Cookie',): 320, ('Green Tea',): 310, ('Cheese Croissant',): 385, ('Blackberry Tart',): 380, ('Single Espresso',): 327}, 2: {'Apricot Danish', 'Cherry Tart': 256}}

#### Rules (Bakery):

[{Cherry Tart} -> {Apricot Danish}, {Apricot Danish} -> {Cherry Tart}]

4. Use Naive bayes, K-nearest, and Decision tree classification algorithms to build classifiers on any two datasets. Pre-process the datasets using techniques specified in Q2. Compare the Accuracy, Precision, Recall and F1 measure reported for each dataset using the abovementioned classifiers under the following situations: i. Using Holdout method (Random sampling):

- a) Training set = 80% Test set = 20%
- b) Training set = 66.6% (2/3rd of total), Test set = 33.3% ii. Using Cross-Validation:
  - a) 10-fold
  - b) 5-fold

```
In [1]: import pandas as pd
import numpy as np
from sklearn.datasets import load_iris, load_wine
from sklearn.model_selection import train_test_split, cross_validate, KFold
from sklearn.preprocessing import StandardScaler
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
In [2]: iris = load_iris()
X1 = pd.DataFrame(iris.data, columns=iris.feature_names)
y1 = iris.target

wine = load_wine()
X2 = pd.DataFrame(wine.data, columns=wine.feature_names)
y2 = wine.target
```

```
In [3]: scaler = StandardScaler()
X1_scaled = scaler.fit_transform(X1)
X2_scaled = scaler.fit_transform(X2)

models = {
    "Naive Bayes": GaussianNB(),
    "KNN": KNeighborsClassifier(n_neighbors=5),
    "Decision Tree": DecisionTreeClassifier(random_state=42)
}

def evaluate_model(model, X_train, X_test, y_train, y_test):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    return {
        "Accuracy": accuracy_score(y_test, y_pred),
        "Precision": precision_score(y_test, y_pred, average='weighted'),
        "Recall": recall_score(y_test, y_pred, average='weighted'),
```

```

    "F1-score": f1_score(y_test, y_pred, average='weighted')
}

```

In [4]: # Holdout Validation

```

holdout_results = []

for dataset_name, (X, y) in {"Iris": (X1_scaled, y1), "Wine": (X2_scaled, y2)}.items():
    for split, test_size in {"80/20": 0.2, "66.6/33.3": 0.333}.items():
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size,
                                                            random_state=42)

        for model_name, model in models.items():
            scores = evaluate_model(model, X_train, X_test, y_train, y_test)
            holdout_results.append({
                "Dataset": dataset_name,
                "Split": split,
                "Model": model_name,
                **scores
            })

holdout_df = pd.DataFrame(holdout_results)
print("\n=== Holdout Method Results ===")
print(holdout_df)

```

=== Holdout Method Results ===

|    | Dataset | Split     | Model         | Accuracy | Precision | Recall   | F1-score |
|----|---------|-----------|---------------|----------|-----------|----------|----------|
| 0  | Iris    | 80/20     | Naive Bayes   | 0.966667 | 0.969697  | 0.966667 | 0.966583 |
| 1  | Iris    | 80/20     | KNN           | 0.933333 | 0.944444  | 0.933333 | 0.932660 |
| 2  | Iris    | 80/20     | Decision Tree | 0.933333 | 0.933333  | 0.933333 | 0.933333 |
| 3  | Iris    | 66.6/33.3 | Naive Bayes   | 0.920000 | 0.923649  | 0.920000 | 0.919722 |
| 4  | Iris    | 66.6/33.3 | KNN           | 0.920000 | 0.935238  | 0.920000 | 0.918877 |
| 5  | Iris    | 66.6/33.3 | Decision Tree | 0.940000 | 0.949000  | 0.940000 | 0.939529 |
| 6  | Wine    | 80/20     | Naive Bayes   | 0.972222 | 0.974359  | 0.972222 | 0.972263 |
| 7  | Wine    | 80/20     | KNN           | 0.972222 | 0.974747  | 0.972222 | 0.972369 |
| 8  | Wine    | 80/20     | Decision Tree | 0.944444 | 0.951389  | 0.944444 | 0.944961 |
| 9  | Wine    | 66.6/33.3 | Naive Bayes   | 0.966667 | 0.969231  | 0.966667 | 0.966456 |
| 10 | Wine    | 66.6/33.3 | KNN           | 0.950000 | 0.957895  | 0.950000 | 0.950476 |
| 11 | Wine    | 66.6/33.3 | Decision Tree | 0.983333 | 0.984000  | 0.983333 | 0.983290 |

In [5]: # Cross-Validation

```

cv_results = []

for dataset_name, (X, y) in {"Iris": (X1_scaled, y1), "Wine": (X2_scaled, y2)}.items():
    for folds in [5, 10]:
        kf = KFold(n_splits=folds, shuffle=True, random_state=42)

        for model_name, model in models.items():
            scores = cross_validate(model, X, y, cv=kf,
                                   scoring=['accuracy', 'precision_weighted', 'recall_weighted'])

            cv_results.append({
                "Dataset": dataset_name,
                "CV_Folds": folds,
                "Model": model_name,
                "Accuracy": np.mean(scores['test_accuracy']),
            })

```

```

        "Precision": np.mean(scores['test_precision_weighted']),
        "Recall": np.mean(scores['test_recall_weighted']),
        "F1-score": np.mean(scores['test_f1_weighted'])
    })

cv_df = pd.DataFrame(cv_results)
print("\n=== Cross-Validation Results ===")
print(cv_df)

```

=== Cross-Validation Results ===

|    | Dataset | CV_Folds | Model         | Accuracy | Precision | Recall   | F1-score |
|----|---------|----------|---------------|----------|-----------|----------|----------|
| 0  | Iris    | 5        | Naive Bayes   | 0.960000 | 0.963879  | 0.960000 | 0.959916 |
| 1  | Iris    | 5        | KNN           | 0.960000 | 0.962187  | 0.960000 | 0.959797 |
| 2  | Iris    | 5        | Decision Tree | 0.953333 | 0.958657  | 0.953333 | 0.953185 |
| 3  | Iris    | 10       | Naive Bayes   | 0.960000 | 0.967619  | 0.960000 | 0.960287 |
| 4  | Iris    | 10       | KNN           | 0.953333 | 0.960254  | 0.953333 | 0.952893 |
| 5  | Iris    | 10       | Decision Tree | 0.940000 | 0.948571  | 0.940000 | 0.940219 |
| 6  | Wine    | 5        | Naive Bayes   | 0.983016 | 0.983844  | 0.983016 | 0.982961 |
| 7  | Wine    | 5        | KNN           | 0.949365 | 0.956877  | 0.949365 | 0.949228 |
| 8  | Wine    | 5        | Decision Tree | 0.876349 | 0.882125  | 0.876349 | 0.876774 |
| 9  | Wine    | 10       | Naive Bayes   | 0.977778 | 0.981411  | 0.977778 | 0.977469 |
| 10 | Wine    | 10       | KNN           | 0.971895 | 0.977178  | 0.971895 | 0.972372 |
| 11 | Wine    | 10       | Decision Tree | 0.898693 | 0.906694  | 0.898693 | 0.897523 |

```

In [6]: # Comparison of Holdout vs Cross-Validation
print("\n===== FINAL COMPARISON =====")
print(pd.concat([holdout_df, cv_df], ignore_index=True))

```

===== FINAL COMPARISON =====

|    | Dataset | Split     | Model         | Accuracy | Precision | Recall   | F1-score | \ |
|----|---------|-----------|---------------|----------|-----------|----------|----------|---|
| 0  | Iris    | 80/20     | Naive Bayes   | 0.966667 | 0.969697  | 0.966667 | 0.966583 |   |
| 1  | Iris    | 80/20     | KNN           | 0.933333 | 0.944444  | 0.933333 | 0.932660 |   |
| 2  | Iris    | 80/20     | Decision Tree | 0.933333 | 0.933333  | 0.933333 | 0.933333 |   |
| 3  | Iris    | 66.6/33.3 | Naive Bayes   | 0.920000 | 0.923649  | 0.920000 | 0.919722 |   |
| 4  | Iris    | 66.6/33.3 | KNN           | 0.920000 | 0.935238  | 0.920000 | 0.918877 |   |
| 5  | Iris    | 66.6/33.3 | Decision Tree | 0.940000 | 0.949000  | 0.940000 | 0.939529 |   |
| 6  | Wine    | 80/20     | Naive Bayes   | 0.972222 | 0.974359  | 0.972222 | 0.972263 |   |
| 7  | Wine    | 80/20     | KNN           | 0.972222 | 0.974747  | 0.972222 | 0.972369 |   |
| 8  | Wine    | 80/20     | Decision Tree | 0.944444 | 0.951389  | 0.944444 | 0.944961 |   |
| 9  | Wine    | 66.6/33.3 | Naive Bayes   | 0.966667 | 0.969231  | 0.966667 | 0.966456 |   |
| 10 | Wine    | 66.6/33.3 | KNN           | 0.950000 | 0.957895  | 0.950000 | 0.950476 |   |
| 11 | Wine    | 66.6/33.3 | Decision Tree | 0.983333 | 0.984000  | 0.983333 | 0.983290 |   |
| 12 | Iris    | NaN       | Naive Bayes   | 0.960000 | 0.963879  | 0.960000 | 0.959916 |   |
| 13 | Iris    | NaN       | KNN           | 0.960000 | 0.962187  | 0.960000 | 0.959797 |   |
| 14 | Iris    | NaN       | Decision Tree | 0.953333 | 0.958657  | 0.953333 | 0.953185 |   |
| 15 | Iris    | NaN       | Naive Bayes   | 0.960000 | 0.967619  | 0.960000 | 0.960287 |   |
| 16 | Iris    | NaN       | KNN           | 0.953333 | 0.960254  | 0.953333 | 0.952893 |   |
| 17 | Iris    | NaN       | Decision Tree | 0.940000 | 0.948571  | 0.940000 | 0.940219 |   |
| 18 | Wine    | NaN       | Naive Bayes   | 0.983016 | 0.983844  | 0.983016 | 0.982961 |   |
| 19 | Wine    | NaN       | KNN           | 0.949365 | 0.956877  | 0.949365 | 0.949228 |   |
| 20 | Wine    | NaN       | Decision Tree | 0.876349 | 0.882125  | 0.876349 | 0.876774 |   |
| 21 | Wine    | NaN       | Naive Bayes   | 0.977778 | 0.981411  | 0.977778 | 0.977469 |   |
| 22 | Wine    | NaN       | KNN           | 0.971895 | 0.977178  | 0.971895 | 0.972372 |   |
| 23 | Wine    | NaN       | Decision Tree | 0.898693 | 0.906694  | 0.898693 | 0.897523 |   |

CV\_Folds

|    |      |
|----|------|
| 0  | NaN  |
| 1  | NaN  |
| 2  | NaN  |
| 3  | NaN  |
| 4  | NaN  |
| 5  | NaN  |
| 6  | NaN  |
| 7  | NaN  |
| 8  | NaN  |
| 9  | NaN  |
| 10 | NaN  |
| 11 | NaN  |
| 12 | 5.0  |
| 13 | 5.0  |
| 14 | 5.0  |
| 15 | 10.0 |
| 16 | 10.0 |
| 17 | 10.0 |
| 18 | 5.0  |
| 19 | 5.0  |
| 20 | 5.0  |
| 21 | 10.0 |
| 22 | 10.0 |
| 23 | 10.0 |

```
In [7]: # Visualize Decision Tree structures for Iris and Wine
from sklearn.tree import plot_tree
import matplotlib.pyplot as plt
```

```


# Train a shallow tree for readability (on original features)
dt_iris = DecisionTreeClassifier(random_state=42, max_depth=5)
dt_iris.fit(X1, y1)

plt.figure(figsize=(12, 8))
plot_tree(dt_iris)
plt.title("Decision Tree (Iris)")
plt.show()


dt_wine = DecisionTreeClassifier(random_state=42, max_depth=5)
dt_wine.fit(X2, y2)

plt.figure(figsize=(14, 10))
plot_tree(dt_wine)
plt.title("Decision Tree (Wine)")
plt.show()

```



# Decision Tree (Wine)



5. Apply simple K-means algorithm for clustering any dataset. Compare the performance of clusters by varying the algorithm parameters. For a given set of parameters, plot a line graph depicting MSE obtained after each iteration.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
```

```
In [2]: # Load Iris dataset
iris = load_iris()
X = iris.data
y = iris.target

# Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

print("Iris Dataset Shape:", X.shape)
print("Features:", iris.feature_names)
print("Target classes:", iris.target_names)
```

Iris Dataset Shape: (150, 4)

Features: ['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']

Target classes: ['setosa' 'versicolor' 'virginica']

```
In [3]: def kmeans(X, n_clusters=3, max_iter=100, random_state=42):
    """
    Simple K-Means implementation that tracks MSE per iteration
    Returns: centroids, labels, mse_history
    """
    np.random.seed(random_state)
    n_samples, n_features = X.shape

    # Initialize centroids randomly from the data points
    random_indices = np.random.choice(n_samples, n_clusters, replace=False)
    centroids = X[random_indices].copy()

    mse_history = []
    labels = None

    for iteration in range(max_iter):
        # Assign each point to the nearest centroid
        distances = np.zeros((n_samples, n_clusters))
        for k in range(n_clusters):
            distances[:, k] = np.linalg.norm(X - centroids[k], axis=1)
        labels = np.argmin(distances, axis=1)

        # Calculate MSE for this iteration
        total_error = 0
```



```

    for k in range(n_clusters):
        cluster_points = X[labels == k]
        if len(cluster_points) > 0:
            total_error += np.sum((cluster_points - centroids[k]) ** 2)
    mse = total_error / n_samples
    mse_history.append(mse)

    # Update centroids
    new_centroids = np.zeros((n_clusters, n_features))
    for k in range(n_clusters):
        cluster_points = X[labels == k]
        if len(cluster_points) > 0:
            new_centroids[k] = cluster_points.mean(axis=0)
        else:
            new_centroids[k] = centroids[k]

    # Check for convergence
    if np.allclose(centroids, new_centroids):
        print(f"Converged at iteration {iteration + 1}")
        break

    centroids = new_centroids

    return centroids, labels, mse_history

```

```

In [4]: # Apply K-Means with k=3 (since Iris has 3 species)
centroids, labels, mse_history = kmeans(X_scaled, n_clusters=3, max_iter=100, random_state=42)

print(f"Final MSE: {mse_history[-1]:.4f}")
print(f"Number of iterations: {len(mse_history)}")

```

```


Converged at iteration 7
Final MSE: 0.9336
Number of iterations: 7

```

```

In [5]: # Plot MSE vs Iteration
plt.figure(figsize=(10, 6))
plt.plot(range(1, len(mse_history) + 1), mse_history, 'b-o', linewidth=2, markersize=10)
plt.xlabel('Iteration', fontsize=12)
plt.ylabel('Mean Squared Error (MSE)', fontsize=12)
plt.title('K-Means: MSE vs Iteration (k=3)', fontsize=14)
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

```




```
In [6]: k_values = [2, 3, 4, 5, 6, 7, 8]
final_mse_values = []
mse_histories = {}

for k in k_values:
    _, _, hist = kmeans(X_scaled, n_clusters=k, max_iter=100, random_state=42)
    final_mse_values.append(hist[-1])
    mse_histories[k] = hist
    print(f"k={k}: Final MSE = {hist[-1]:.4f}, Iterations = {len(hist)}")
```


```
Converged at iteration 3
k=2: Final MSE = 1.4824, Iterations = 3
Converged at iteration 7
k=3: Final MSE = 0.9336, Iterations = 7
Converged at iteration 10
k=4: Final MSE = 0.7637, Iterations = 10
Converged at iteration 7
k=5: Final MSE = 0.6908, Iterations = 7
Converged at iteration 9
k=6: Final MSE = 0.5392, Iterations = 9
Converged at iteration 7
k=7: Final MSE = 0.4811, Iterations = 7
Converged at iteration 8
k=8: Final MSE = 0.4486, Iterations = 8
```

```
In [7]: # Elbow Method: Final MSE vs k
plt.figure(figsize=(10, 6))
plt.plot(k_values, final_mse_values, 'r-o', linewidth=2, markersize=10)
plt.xlabel('Number of Clusters (k)', fontsize=12)
plt.ylabel('Final MSE', fontsize=12)
plt.title('Elbow Method: Final MSE vs Number of Clusters', fontsize=14)
plt.xticks(k_values)
plt.grid(True, alpha=0.3)
```

```
plt.tight_layout()
plt.show()
```



```
In [8]: # MSE per iteration for different k values
plt.figure(figsize=(12, 6))
for k in [2, 3, 4, 5]:
    plt.plot(range(1, len(mse_histories[k]) + 1), mse_histories[k], '-o', label=f'k={k}')
plt.xlabel('Iteration', fontsize=12)
plt.ylabel('Mean Squared Error (MSE)', fontsize=12)
plt.title('MSE vs Iteration for Different k Values', fontsize=14)
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```




```
In [9]: random_states = [0, 10, 42, 100, 200]
init_mse_histories = {}

plt.figure(figsize=(12, 6))
for rs in random_states:
    _, _, hist = kmeans(X_scaled, n_clusters=3, max_iter=100, random_state=rs)
    init_mse_histories[rs] = hist
    print(f"Random State {rs}: Final MSE = {hist[-1]:.4f}, Iterations = {len(hist)}")
    plt.plot(range(1, len(hist) + 1), hist, '-o', label=f'seed={rs}', markersize=6)

plt.xlabel('Iteration', fontsize=12)
plt.ylabel('Mean Squared Error (MSE)', fontsize=12)
plt.title('MSE vs Iteration for Different Random Initializations (k=3)', fontsize=12)
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```

```
Converged at iteration 6
Random State 0: Final MSE = 0.9393, Iterations = 6
Converged at iteration 7
Random State 10: Final MSE = 0.9393, Iterations = 7
Converged at iteration 7
Random State 42: Final MSE = 0.9336, Iterations = 7
Converged at iteration 8
Random State 100: Final MSE = 0.9336, Iterations = 8
Converged at iteration 7
Random State 200: Final MSE = 0.9336, Iterations = 7
```



```
In [10]: # Final clustering with k=3
final_centroids, final_labels, _ = kmeans(X_scaled, n_clusters=3, max_iter=100, ran

# Plot clusters using first two features
plt.figure(figsize=(14, 5))

plt.subplot(1, 2, 1)
scatter = plt.scatter(X_scaled[:, 0], X_scaled[:, 1], c=final_labels, cmap='viridis')
plt.scatter(final_centroids[:, 0], final_centroids[:, 1],
```

```


c='red', marker='X', s=200, edgecolors='black', linewidths=2, label='Ce
plt.xlabel('Sepal Length (scaled)', fontsize=12)
plt.ylabel('Sepal Width (scaled)', fontsize=12)
plt.title('K-Means Clustering (k=3)', fontsize=14)
plt.legend()
plt.colorbar(scatter, label='Cluster')

plt.subplot(1, 2, 2)
scatter = plt.scatter(X_scaled[:, 0], X_scaled[:, 1], c=y, cmap='viridis', alpha=0.
plt.xlabel('Sepal Length (scaled)', fontsize=12)
plt.ylabel('Sepal Width (scaled)', fontsize=12)
plt.title('Actual Iris Species', fontsize=14)
plt.colorbar(scatter, label='Species')

plt.tight_layout()
plt.show()

```

Converged at iteration 7



```

In [11]: # Summary Table
summary_data = {
    'k': k_values,
    'Final MSE': [f"{mse:.4f}" for mse in final_mse_values],
    'Iterations': [len(mse_histories[k]) for k in k_values]
}
summary_df = pd.DataFrame(summary_data)
print("Performance by Number of Clusters:")
print(summary_df.to_string(index=False))

print("\nConclusions:")
print("- As k increases, MSE decreases (more clusters = less error)")
print("- The 'elbow' in the MSE curve suggests optimal k")
print("- Different initializations can lead to different final solutions")
print("- For Iris dataset, k=3 is optimal (matching 3 species)")

```

#### Performance by Number of Clusters:

| k | Final MSE | Iterations |
|---|-----------|------------|
| 2 | 1.4824    | 3          |
| 3 | 0.9336    | 7          |
| 4 | 0.7637    | 10         |
| 5 | 0.6908    | 7          |
| 6 | 0.5392    | 9          |
| 7 | 0.4811    | 7          |
| 8 | 0.4486    | 8          |

#### Conclusions:

- As k increases, MSE decreases (more clusters = less error)
- The 'elbow' in the MSE curve suggests optimal k
- Different initializations can lead to different final solutions
- For Iris dataset, k=3 is optimal (matching 3 species)

## 6. Perform density-based clustering algorithm on a downloaded dataset and evaluate the cluster quality by changing the algorithm's parameters.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import DBSCAN
from sklearn.metrics import silhouette_score, calinski_harabasz_score, davies_bouldin_score
from sklearn.neighbors import NearestNeighbors
```

```
In [2]: # Load and prepare data
df = pd.read_excel('Mall Customers.xlsx')
print("Dataset Shape:", df.shape)
print("Columns:", df.columns.tolist())
df.head()
```

Dataset Shape: (200, 7)

Columns: ['CustomerID', 'Gender', 'Age', 'Education', 'Marital Status', 'Annual Income (k\$)', 'Spending Score (1-100)']

```
Out[2]:
```

|   | CustomerID | Gender | Age | Education   | Marital Status | Annual Income (k\$) | Spending Score (1-100) |
|---|------------|--------|-----|-------------|----------------|---------------------|------------------------|
| 0 | 1          | M      | 19  | High School | Married        | 15                  | 39                     |
| 1 | 2          | M      | 21  | Graduate    | Single         | 15                  | 81                     |
| 2 | 3          | F      | 20  | Graduate    | Married        | 16                  | 6                      |
| 3 | 4          | F      | 23  | High School | Unknown        | 16                  | 77                     |
| 4 | 5          | F      | 31  | Uneducated  | Married        | 17                  | 40                     |

```
In [3]: # Select and scale features
X = df[['Annual Income (k$)', 'Spending Score (1-100)']].values
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
print("Scaled data shape:", X_scaled.shape)
```

Scaled data shape: (200, 2)

```
In [4]: # Helper function to evaluate DBSCAN
def evaluate_dbscan(X, eps, min_samples):
    labels = DBSCAN(eps=eps, min_samples=min_samples).fit_predict(X)
    n_clusters = len(set(labels)) - (1 if -1 in labels else 0)
    n_noise = list(labels).count(-1)

    # Calculate silhouette score (needs at least 2 clusters, excluding noise)
    mask = labels != -1
    silhouette = silhouette_score(X[mask], labels[mask]) if n_clusters >= 2 and sum(mask) > 0 else None
    return n_clusters, n_noise, silhouette, labels
```

```
In [5]: # Evaluate by varying eps (epsilon)
print("Effect of varying eps (min_samples=5):")
print("-" * 45)
for eps in [0.3, 0.4, 0.5, 0.6, 0.7, 0.8]:
    n_clusters, n_noise, sil, _ = evaluate_dbscan(X_scaled, eps, 5)
    sil_str = f"{sil:.3f}" if sil else "N/A"
    print(f"eps={eps}: Clusters={n_clusters}, Noise={n_noise}, Silhouette={sil_str}")
```

Effect of varying eps (min\_samples=5):

```
-----
eps=0.3: Clusters=7, Noise=35, Silhouette=0.524
eps=0.4: Clusters=4, Noise=15, Silhouette=0.478
eps=0.5: Clusters=2, Noise=8, Silhouette=0.388
eps=0.6: Clusters=1, Noise=5, Silhouette=N/A
eps=0.7: Clusters=1, Noise=0, Silhouette=N/A
eps=0.8: Clusters=1, Noise=0, Silhouette=N/A
```

```
In [6]: # Evaluate by varying min_samples
print("Effect of varying min_samples (eps=0.5):")
print("-" * 50)
for min_samples in [3, 4, 5, 6, 7, 10]:
    n_clusters, n_noise, sil, _ = evaluate_dbscan(X_scaled, 0.5, min_samples)
    sil_str = f"{sil:.3f}" if sil else "N/A"
    print(f"min_samples={min_samples}: Clusters={n_clusters}, Noise={n_noise}, Silh
```

Effect of varying min\_samples (eps=0.5):

```
-----
min_samples=3: Clusters=2, Noise=7, Silhouette=0.389
min_samples=4: Clusters=2, Noise=8, Silhouette=0.388
min_samples=5: Clusters=2, Noise=8, Silhouette=0.388
min_samples=6: Clusters=2, Noise=11, Silhouette=0.401
min_samples=7: Clusters=2, Noise=12, Silhouette=0.394
min_samples=10: Clusters=4, Noise=21, Silhouette=0.511
```

```
In [7]: # Visualize final clustering with chosen parameters
eps, min_samples = 0.5, 5
n_clusters, n_noise, sil, labels = evaluate_dbscan(X_scaled, eps, min_samples)

plt.figure(figsize=(10, 6))
plt.scatter(X[:, 0], X[:, 1], c=labels, cmap='viridis', alpha=0.7, s=50)
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.title(f'DBSCAN Clustering (eps={eps}, min_samples={min_samples})\nClusters: {n_clusters}, Noise: {n_noise}')
plt.colorbar(label='Cluster (-1 = Noise)')
plt.tight_layout()
plt.show()
```



DBSCAN Clustering (eps=0.5, min\_samples=5)  
Clusters: 2, Noise: 8, Silhouette: 0.388

